

C3: More on Association Rules: Data Stream, FP-Tree, Vertical Mining

October 2, 2018

Ming-Syan Chen

Agenda for Classes 2018 (Might be revised as we progress)

- Class 1 – (9/11) Overview of data mining
- Class 2 – (9/18) [R \(I\)](#), [Wush Wu](#)
- Class 3 – (9/25) Association, Apriori and its related issues
- Class 4 – (10/2) [Data stream mining](#), [FP Tree](#), [Vertical mining \(announcing HW1\)](#)
- Class 5 – (10/9) Classification: decision tree, [GPGPU](#)
- Class 6 – (10/16) Description of Data, Project announcement (announcing HW2)
- Class 7 – (10/23) [R \(II\)](#), [Wush Wu](#)

Tentative Class Agenda (cont'd)

- Class 8 – (10/30) Data exploration, more on decision trees, rule-based classifiers
- Class 9 – (11/6) [Scikit learn](#), [LibSVM](#), [Preparation for HW3 and HW4](#)
- Class 10 – (11/13) KNN, Bays, Neural network, Concept of SVM
- Class 11 – (11/20) Abstract presentation, SVM, Clustering, K-means, PAM
- Class 12 – (11/27) More on clustering; Sequential pattern mining;
- Class 13 – (12/4) Web mining, PageRank, etc.

Tentative Class Agenda (cont'd)

- Class 14– (12/11) [When Database and Data Mining Meet, Prof. Mingling Lo](#)
- Class 15 – (12/18) Project presentation I
- Class 16 – (12/25) Project presentation II
(Final Exam according to Univ. Schedule)
(Project due 1/24/2019)
- Happy New Year!

Data Stream: Too fast and too huge to capture

Related Papers

- W.-G. Teng, M.-S. Chen and P. S. Yu, "A Regression-Based Temporal Pattern Mining Scheme for Data Streams," *Proc. of the 29th Intern'l Conf. on Very Large Data Bases (VLDB-2003)*, September 9-12, 2003.
- J. Han, J. Pei, Y. Yin, R. Mao, "Mining Frequent Patterns without Candidate Generation: A Frequent-Pattern Tree Approach" DMKD, Vol. 8, No. 1, Jan. 2004. (SIGMOD 2000)
- M. J. Zaki, Scalable Algorithms for Association Mining, IEEE Transactions on Knowledge and Data Engineering 12 (3), 372-390, 2000

Data Streams

- Traditional DBMS – data stored in **finite, persistent** data sets
- Emerging Applications – data input as continuous, ordered data streams
 - Network monitoring and traffic
 - Telecom call detail records (CDR)
 - ATM operations in banks
 - Sensor networks
 - Web logs and click-streams
 - Transactions in retail chains
 - Mobile computing
 - One of the most important issues in data mining

Data Streams (cont'd)

- Definition
 - **Continuous, unbounded, rapid, time-varying** streams of data elements
- Application Characteristics
 - Massive volumes of data (can be several terabytes)
 - Records arrive at a rapid rate
- Goal
 - Mine patterns, process queries and compute statistics on data streams in real-time

Online Transaction Flows

- Example data of a market-basket application
 - In the form of <TxTime, CustomerID, Itemset>

Customer ID	1	(c)	(i)	(g)
	2	(a,b,c)	(c,g)	(d,f,g)
	3	(c)	(c,e,g)	(g)
	4	(c)	(d,g)	(i)
	5	(i)	(c)	(g)


time 0 1 2 3 4 5 now

Support Framework for Temporal Patterns

- Goal: to mine all frequent temporal patterns satisfying the time constraint
- Occurrence frequency (or support) of patterns is the metric commonly used
 - However, the definition of support varies from one application to another
- With the sliding window model, the support of a temporal pattern X is generally (newly) defined as

$$\text{Sup}_t(X) = \frac{\# \text{ of customers having pattern X at time t}}{\# \text{ of all customers}}$$

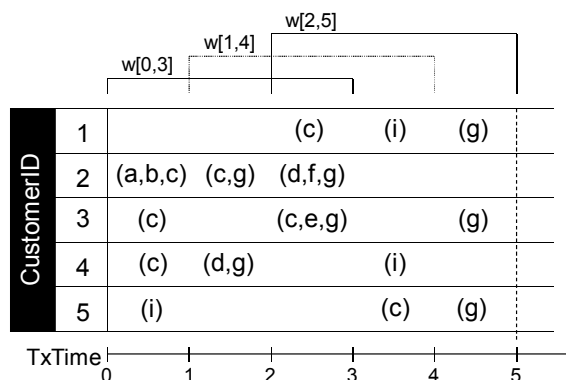
Sliding Window Model



- A fixed-width window is utilized so that only data elements within the current time window are taken into consideration
- Why?
 - Approximation technique for bounded memory
 - Natural in applications (emphasizes recent data)
 - Well-specified and deterministic semantics


Example of the Support Counting

- Consider the itemset $\{c, g\}$ ($wSize=3$)



Example of Support Counting (cont'd)

- Support variations of {c, g}



One Scan for Statistics Collection

- Levelwise generation of patterns
 - In the 1st window, only singleton items are counted
 - Frequent item(set)s are retained and joined to form longer candidate itemsets
 - Candidate itemsets are counted in the next window, while infrequent ones are discarded

Example of One Scan Support Counting ({c,g} not the same as shown before)

- (accumulated supports)/(# of recorded windows)

t=1	
{c}	0.6 / 1

t=2	
{c}	1.2 / 2
{g}	0.4 / 1

t=3	
{c}	2 / 3
{d}	0.4 / 1
{g}	1 / 2
{c, g}	0.6 / 1

t=4	
{c}	2.8 / 4
{d}	0.8 / 2
{g}	1.6 / 3
{i}	0.4 / 1
{c, g}	1 / 2
{d, g}	0.4 / 1

t=5	
{c}	3.4 / 5
{g}	2.4 / 4
{i}	0.8 / 2
{c, g}	1.6 / 3



MinSup=0.4

Delayed Pattern Recognition

- E.g., for {c} t=2, add four 0.2's (1-4; to t=3), add four 0.2's (1,2,3,5; to t=4), add three 0.2's (1,3,5; to t=5)
- Delayed pattern recognition occurs, e.g.,
 - Actually, {c, g} is frequent at t=2
 - But since items c & g are not both frequent until t=2, {c, g} is generated & counted at t=3
 - Only could transient patterns be neglected
 - Patterns with supports near the threshold are examined and identified if so qualified

Regression-Based Analysis

- To maintain support variations of frequent patterns with lower space cost
 - Make the series averaged & find a fit line



$$y = -0.0333x + 0.6778$$

Construct the Fit Line


- To construct the fit line $f=A+Bt$, we have
 - $B = S_{tf}/S_{tt}$, and
 - $A = (\sum f)/n - B(\sum t)/n$,
 where
 - $S_{tt} = \sum t^2 - (\sum t)^2/n$
 - $S_{ff} = \sum f^2 - (\sum f)^2/n$
 - $S_{tf} = \sum tf - (\sum t)(\sum f)/n$
 - Note: n (# of points), $(\sum t)$ and $(\sum t^2)$ can be calculated from t_s and t_{now}

Compact ATF Representation

- Accumulated Time & Frequency form
 - $(t_s, \sum tf, \sum f, \sum f^2)$ can losslessly represent the fit line
 - t_s : starting time
 - t : time index
 - f : occurrence frequency of the pattern at t
 - Easy to update this form by accumulating corresponding (product) entries

Piecewise Linear Representation


- Not only a single fit line but a set of fit lines are used to represent the original series
- Use an ATF list to maintain multiple entries of successive ATF entries



Problem!
Unbounded memory
space is required as
time advances!

Segment Relaxation

- To make memory space bounded
- Temporal granularity is introduced
 - People are interested in recent changes at a fine scale, but old changes at a coarse scale



FP-Tree:
Another way for association rule
mining (in addition to Apriori)

Related Papers

- W.-G. Teng, M.-S. Chen and P. S. Yu, "A Regression-Based Temporal Pattern Mining Scheme for Data Streams," *Proc. of the 29th Intern'l Conf. on Very Large Data Bases (VLDB-2003)*, September 9-12, 2003.
- J. Han, J. Pei, Y. Yin, R. Mao, "Mining Frequent Patterns without Candidate Generation: A Frequent-Pattern Tree Approach" *DMKD*, Vol. 8, No. 1, Jan. 2004. (SIGMOD 2000)
- M. J. Zaki, Scalable Algorithms for Association Mining, *IEEE Transactions on Knowledge and Data Engineering* 12 (3), 372-390, 2000

Apriori-like Algorithm

- Essential idea
 - Generate the set of candidate patterns of length $(k+1)$ from the set of frequent patterns of length k .
 - Check the corresponding occurrence frequencies of the candidate patterns in the database.

What does Apriori-like algorithm suffer from?

- In situations with many frequent patterns, long patterns, or quite low minimum support threshold, it is costly.
- It is tedious to repeatedly scan the database and check a large set of candidates by patterns matching.

FP-Tree: To avoid generating candidate patterns

- A highly compact data structure: frequent pattern tree
- An FP-tree-based pattern fragment growth mining method
- The paper on SIGMOD 2000 has spawned many subsequent studies
- Note however when the database is large, it becomes not feasible to put the whole tree into the memory (i.e., FP-tree is usually good for small db only)


Example 1 (minsup=3)

Tx ID	Items Bought	(ordered) Frequent Items
100	f,a,c,d,g,i,m,p	f,c,a,m,p
200	a,b,c,f,l,m,o	f,c,a,b,m
300	b,f,h,j,o	f,b
400	b,c,k,s,p	c,b,p
500	a,f,c,e,l,p,m,n	f,c,a,m,p

List of frequent items:
(f:4),(c:4),(a:3),(b:3),(m:3),(p:3)

FP-tree

Header table



Frequent pattern tree (FP-tree)

- 3 parts
 - One root labeled as “null”
 - A set of item prefix subtrees
 - Frequent item header table

FP-tree(cont’)

- Each node in the prefix subtree consists of
 - Item-name
 - Count
 - Node-link
- Each entry in the frequent-item header table consists of
 - Item-name
 - Head of node-link

FP-tree construction: step 1

- Scan the transaction database DB once (**the first time**), and derives a list of frequent items.
- Sort frequent items in frequency descending order.
- This ordering is important since each path of a tree will follow this order

List of frequent items:

(f:4),(c:4),(a:3),(b:3),(m:3),(p:3)

FP-tree construction: step 2

- Create a root of a tree, label with “null”
- Scan the database **the second time**.


The scan of the first tx leads to the construction of the first branch of the tree.

FP-tree construction: step 2(cont')

Scan of 1st Transaction : f,a,c,d,g,i,m,p

The 1st branch of the tree

<(f:1),(c:1),(a:1),(m:1),(p:1)>



FP-tree construction: step 2(cont')


Scan of 2nd Transaction :

a,b,c,f,l,m,o

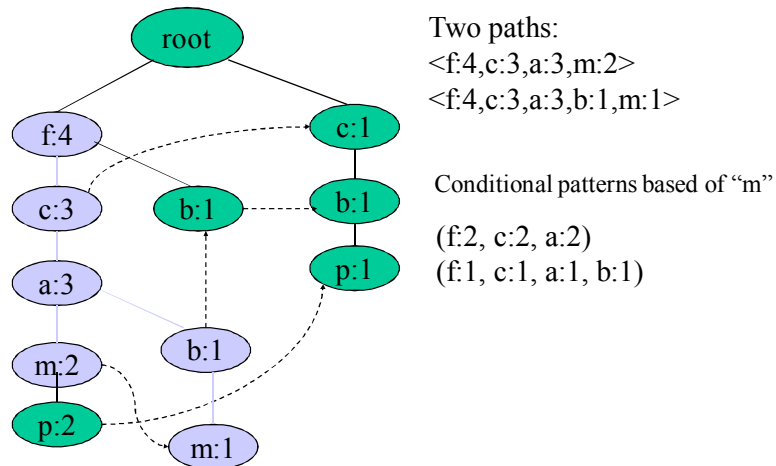
(-> f, c, a, b, m)

Two new nodes:


(b:1) (m:1)




Mining Process: for node “m”



Conditional FP-tree of “m”



Mining m's conditional FP-tree



Mining m's conditional FP-tree

- The whole set of frequent patterns
 $\{(m:3), (am:3), (cm:3), (fm:3), (cam:3), (fam:3), (fcam:3), (fcm:3)\}$

Note: We perform the identification of frequent patterns from the leaf part of the tree first.

Remarks

- Association rule mining starts from 1993 (term used) and has been well received by the community.
- Both Apriori and FP-Tree approaches have spawned many subsequent studies (and variations)

Vertical Mining

Related Papers

- W.-G. Teng, M.-S. Chen and P. S. Yu, "A Regression-Based Temporal Pattern Mining Scheme for Data Streams," *Proc. of the 29th Intern'l Conf. on Very Large Data Bases (VLDB-2003)*, September 9-12, 2003.
- J. Han, J. Pei, Y. Yin, R. Mao, "Mining Frequent Patterns without Candidate Generation: A Frequent-Pattern Tree Approach" *DMKD*, Vol. 8, No. 1, Jan. 2004. (SIGMOD 2000)
- [M. J. Zaki, Scalable Algorithms for Association Mining, IEEE Transactions on Knowledge and Data Engineering 12 \(3\), 372-390, 2000](#)

Frequent Itemset Mining (FIP)

- Notation
 - Items I , Transaction Database T
 - Itemset: combination of distinct items (e.g. AD, CW ...)
 - $\sigma(X)$, support of an itemset, is the number of transactions in which X occur
- Find all itemsets with support $> min_sup$

DATABASE		ALL FREQUENT ITEMSETS MINIMUM SUPPORT = 50%	
Transaction	Items	Support	Itemsets
1	A C T W	100% (6)	C
2	C D W	83% (5)	W, CW
3	A C T W	67% (4)	A, D, T, AC, AW CD, CT, ACW
4	A C D W	50% (3)	AT, DW, TW, ACT, ATW CDW, CTW, ACTW
5	A C D T W		
6	C D T		

Database Representation

- Horizontal items

Transaction ID	Items
1	A C T W
2	C D W
3	A C T W
4	A C D W
5	A C D T W
6	C D T

Drawback: Need to traverse overall dataset to count support

Database Representation(cont'd)

- Vertical Transaction ID set (Tidset)

Itemset	A	C	D	T	W
Tidset	1 3 4 5	1 2 3 4 5 6	2 4 5 6	1 3 5 6	1 2 3 4 5

Efficient for counting support !

Vertical Mining

Eclat[M. J. Zaki et al., TKDE 2000]

- 1. Intersect each two Tidset with the same itemset prefix
- 2. Count support and prune infrequent itemset
- 3. Continue until no new frequent itemset

min_sup: 3

